

FraudDetection_with_Preprocessing

November 3, 2024

Name: Hamid El Messaoudi

Date: 11/3/2024

Program: Complete Data Preprocessing Workflow for Fraud Detection

```
[55]: # Importing the basic libraries we need for data manipulation and analysis
import pandas as pd
import numpy as np
```

```
[56]: # Installing necessary packages (only run if not installed)
!python -m pip install --upgrade pip
!pip install sklearn_pandas xgboost
```

Requirement already satisfied: pip in

c:\users\hamid\appdata\local\programs\python\python312\lib\site-packages
(24.3.1)

Requirement already satisfied: sklearn_pandas in

c:\users\hamid\appdata\local\programs\python\python312\lib\site-packages (2.2.0)

Requirement already satisfied: xgboost in

c:\users\hamid\appdata\local\programs\python\python312\lib\site-packages (2.1.2)

Requirement already satisfied: scikit-learn>=0.23.0 in

c:\users\hamid\appdata\local\programs\python\python312\lib\site-packages (from
sklearn_pandas) (1.4.0)

Requirement already satisfied: scipy>=1.5.1 in

c:\users\hamid\appdata\local\programs\python\python312\lib\site-packages (from
sklearn_pandas) (1.12.0)

Requirement already satisfied: pandas>=1.1.4 in

c:\users\hamid\appdata\local\programs\python\python312\lib\site-packages (from
sklearn_pandas) (2.2.0)

Requirement already satisfied: numpy>=1.18.1 in

c:\users\hamid\appdata\local\programs\python\python312\lib\site-packages (from
sklearn_pandas) (1.26.3)

Requirement already satisfied: python-dateutil>=2.8.2 in

c:\users\hamid\appdata\local\programs\python\python312\lib\site-packages (from
pandas>=1.1.4->sklearn_pandas) (2.8.2)

Requirement already satisfied: pytz>=2020.1 in

c:\users\hamid\appdata\local\programs\python\python312\lib\site-packages (from
pandas>=1.1.4->sklearn_pandas) (2024.1)

Requirement already satisfied: tzdata>=2022.7 in
c:\users\hamid\appdata\local\programs\python\python312\lib\site-packages (from
pandas>=1.1.4->sklearn_pandas) (2023.4)
Requirement already satisfied: joblib>=1.2.0 in
c:\users\hamid\appdata\local\programs\python\python312\lib\site-packages (from
scikit-learn>=0.23.0->sklearn_pandas) (1.3.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in
c:\users\hamid\appdata\local\programs\python\python312\lib\site-packages (from
scikit-learn>=0.23.0->sklearn_pandas) (3.2.0)
Requirement already satisfied: six>=1.5 in
c:\users\hamid\appdata\local\programs\python\python312\lib\site-packages (from
python-dateutil>=2.8.2->pandas>=1.1.4->sklearn_pandas) (1.16.0)

```
[57]: # Reading the dataset
data = pd.read_csv('insuranceFraud.csv')
```

```
[58]: # Having a look at the data
data.head()
```

```
[58]:  months_as_customer  age  policy_number  policy_bind_date  policy_state  \
0                328   48         521585      10/17/2014          OH
1                228   42         342868       6/27/2006          IN
2                134   29         687698       9/6/2000          OH
3                256   41         227811      5/25/1990          IL
4                228   44         367455       6/6/2014          IL

   policy_csl  policy_deductable  policy_annual_premium  umbrella_limit  \
0    250/500              1000          1406.91              0
1    250/500              2000          1197.22      5000000
2    100/300              2000          1413.14      5000000
3    250/500              2000          1415.74      6000000
4    500/1000             1000          1583.91      6000000

   insured_zip  ... witnesses  police_report_available  total_claim_amount  \
0     466132  ...         2              YES              71610
1     468176  ...         0              ?              5070
2     430632  ...         3              NO              34650
3     608117  ...         2              NO              63400
4     610706  ...         1              NO              6500

   injury_claim  property_claim  vehicle_claim  auto_make  auto_model  auto_year  \
0          6510          13020          52080        Saab         92x         2004
1           780           780           3510    Mercedes         E400         2007
2          7700          3850          23100     Dodge         RAM         2007
3          6340          6340          50720  Chevrolet         Tahoe         2014
4          1300           650           4550    Accura         RSX         2009
```

```

    fraud_reported
0                Y
1                Y
2                N
3                Y
4                N

```

[5 rows x 39 columns]

```

[59]: # Replacing all the "?" values with NaN to make them easier to handle
data = data.replace('?', np.nan)

```

```

[60]: # list of columns not necessary for pfrediction
cols_to_drop=['policy_number', 'policy_bind_date', 'policy_state', 'insured_zip', 'incident_location']

```

```

[61]: # dropping the unnecessary columns
data.drop(columns=cols_to_drop, inplace=True)

```

```

[62]: # checking the data after dropping the columns
data.head()

```

```

[62]:   months_as_customer  age  policy_csl  policy_deductable \
0                328    48    250/500                1000
1                228    42    250/500                2000
2                134    29    100/300                2000
3                256    41    250/500                2000
4                228    44    500/1000               1000

```

```

   policy_annual_premium  umbrella_limit  insured_sex  insured_education_level \
0                1406.91                0        MALE                      MD
1                1197.22            5000000        MALE                      MD
2                1413.14            5000000       FEMALE                      PhD
3                1415.74            6000000       FEMALE                      PhD
4                1583.91            6000000        MALE              Associate

```

```

   insured_occupation  insured_relationship  ...  number_of_vehicles_involved \
0      craft-repair             husband  ...                1
1  machine-op-inspct      other-relative  ...                1
2             sales             own-child  ...                3
3    armed-forces             unmarried  ...                1
4             sales             unmarried  ...                1

```

```

   property_damage  bodily_injuries  witnesses  police_report_available \
0             YES                1            2                YES
1             NaN                0            0                NaN
2             NO                2            3                NO
3             NaN                1            2                NO

```

4	NO	0	1	NO
---	----	---	---	----

	total_claim_amount	injury_claim	property_claim	vehicle_claim	\
0	71610	6510	13020	52080	
1	5070	780	780	3510	
2	34650	7700	3850	23100	
3	63400	6340	6340	50720	
4	6500	1300	650	4550	

	fraud_reported
0	Y
1	Y
2	N
3	Y
4	N

[5 rows x 27 columns]


```
[63]: # checking for missing values
data.isna().sum()
```

```
[63]: months_as_customer      0
age                           0
policy_csl                    0
policy_deductable             0
policy_annual_premium         0
umbrella_limit                0
insured_sex                   0
insured_education_level       0
insured_occupation            0
insured_relationship           0
capital-gains                 0
capital-loss                   0
incident_type                  0
collision_type                 178
incident_severity              0
authorities_contacted         91
incident_hour_of_the_day       0
number_of_vehicles_involved    0
property_damage                360
bodily_injuries                0
witnesses                      0
police_report_available        343
total_claim_amount             0
injury_claim                   0
property_claim                 0
vehicle_claim                  0
```

```
fraud_reported          0
dtype: int64
```

```
[64]: import seaborn as sns
import matplotlib.pyplot as plt
missing = data.isnull().mean() * 100 # percentage
missing = missing[missing > 0]
missing.sort_values(inplace=True)
missing = missing.to_frame()
missing.columns = ['Null_Count']
missing.index.names = ['Col_Name']
missing = missing.reset_index()

sns.set(style='whitegrid', color_codes=True)
sns.barplot(x='Col_Name', y='Null_Count', data=missing)
plt.xticks(rotation=90)
plt.show()
```




```
[65]: # Set the figure size for the correlation heatmap
plt.figure(figsize=(18, 15))

# Calculate the correlation matrix for numerical columns in the data
corr = data.select_dtypes(include=['number']).corr()

# Plot the heatmap of the correlation matrix with annotations and specified
# formatting
sns.heatmap(data=corr, annot=True, fmt='.1g', linewidth=2)


# Display the plot
plt.show()
```



```
[66]: # Checking if any of our numerical features are strongly related to each other_
      ↪(aka, correlated).
      # If two features are super similar, we might not need both.
      # This just gives us a sense of what's related in our data.
      plt.figure(figsize = (18,15))
      corr = data.select_dtypes(include=['number']).corr()
      mask = np.triu(np.ones_like(corr, dtype = bool))

      sns.heatmap(data = corr, mask = mask, annot = True, fmt = '.2g', linewidth = 1)
      plt.show
```

```
[66]: <function matplotlib.pyplot.show(close=None, block=None)>
```



```
[67]: data.drop(columns = ['age', 'total_claim_amount'], inplace = True, axis = 1)
data.head()
```

```
[67]:  months_as_customer  policy_csl  policy_deductable  policy_annual_premium  \
0                328      250/500                1000             1406.91
1                228      250/500                2000             1197.22
2                134      100/300                2000             1413.14
3                256      250/500                2000             1415.74
4                228      500/1000               1000             1583.91

   umbrella_limit  insured_sex  insured_education_level  insured_occupation  \
0                0        MALE                      MD      craft-repair
1          5000000        MALE                      MD  machine-op-inspct
2          5000000       FEMALE                      PhD             sales
3          6000000       FEMALE                      PhD      armed-forces
```


4	6000000	MALE		Associate		sales
---	---------	------	--	-----------	--	-------

	insured_relationship	capital-gains	...	incident_hour_of_the_day	\
0	husband	53300	...	5	
1	other-relative	0	...	8	
2	own-child	35100	...	7	
3	unmarried	48900	...	5	
4	unmarried	66000	...	20	

	number_of_vehicles_involved	property_damage	bodily_injuries	witnesses	\
0	1	YES	1	2	
1	1	NaN	0	0	
2	3	NO	2	3	
3	1	NaN	1	2	
4	1	NO	0	1	

	police_report_available	injury_claim	property_claim	vehicle_claim	\
0	YES	6510	13020	52080	
1	NaN	780	780	3510	
2	NO	7700	3850	23100	
3	NO	6340	6340	50720	
4	NO	1300	650	4550	


	fraud_reported
0	Y
1	Y
2	N
3	Y
4	N

[5 rows x 25 columns]

```
[68]: plt.figure(figsize = (18,15))
corr = data.select_dtypes(include=['number']).corr()
mask = np.triu(np.ones_like(corr,dtype =bool))

sns.heatmap(data =corr, mask = mask, annot = True, fmt ='.2g', linewidth =1)
plt.show
```

```
[68]: <function matplotlib.pyplot.show(close=None, block=None)>
```



```
[69]: # Importing SimpleImputer from sklearn
from sklearn.impute import SimpleImputer

# Creating an imputer instance to fill missing values with the most frequent_
↳ value in each column
imputer = SimpleImputer(strategy='most_frequent')

# Imputing missing values for specific categorical columns
data['collision_type'] = imputer.fit_transform(data[['collision_type']]).ravel()
data['property_damage'] = imputer.fit_transform(data[['property_damage']]).
↳ ravel()
data['police_report_available'] = imputer.
↳ fit_transform(data[['police_report_available']]).ravel()
```

```
[70]: # As the columns which have missing values, they are only categorical, we'll
      ↪ use the categorical imputer
      # Importing the categorical imputer
      from sklearn.impute import SimpleImputer

      # Create an imputer instance for categorical columns
      imputer = SimpleImputer(strategy='most_frequent')

[71]: # Imputing missing values for each categorical column
      data['collision_type'] = imputer.fit_transform(data[['collision_type']]).ravel()
      data['property_damage'] = imputer.fit_transform(data[['property_damage']]).
      ↪ ravel()
      data['police_report_available'] = imputer.
      ↪ fit_transform(data[['police_report_available']]).ravel()

[72]: # Extracting the categorical columns
      cat_df = data.select_dtypes(include=['object']).copy()

[73]: cat_df.columns

[73]: Index(['policy_cs1', 'insured_sex', 'insured_education_level',
          'insured_occupation', 'insured_relationship', 'incident_type',
          'collision_type', 'incident_severity', 'authorities_contacted',
          'property_damage', 'police_report_available', 'fraud_reported'],
          dtype='object')

[74]: cat_df.head()

[74]: policy_cs1 insured_sex insured_education_level insured_occupation \
0      250/500      MALE      MD      craft-repair
1      250/500      MALE      MD  machine-op-inspct
2      100/300    FEMALE      PhD      sales
3      250/500    FEMALE      PhD  armed-forces
4      500/1000     MALE      Associate      sales

      insured_relationship      incident_type collision_type \
0      husband  Single Vehicle Collision  Side Collision
1  other-relative      Vehicle Theft  Rear Collision
2      own-child  Multi-vehicle Collision  Rear Collision
3      unmarried  Single Vehicle Collision  Front Collision
4      unmarried      Vehicle Theft  Rear Collision

      incident_severity authorities_contacted property_damage \
0      Major Damage      Police      YES
1      Minor Damage      Police      NO
2      Minor Damage      Police      NO
3      Major Damage      Police      NO
```

4	Minor Damage	NaN	NO
---	--------------	-----	----

	police_report_available	fraud_reported
0	YES	Y
1	NO	Y
2	NO	N
3	NO	Y
4	NO	N

```
[75]: cat_df.columns
```

```
[75]: Index(['policy_cs1', 'insured_sex', 'insured_education_level',
          'insured_occupation', 'insured_relationship', 'incident_type',
          'collision_type', 'incident_severity', 'authorities_contacted',
          'property_damage', 'police_report_available', 'fraud_reported'],
          dtype='object')
```

```
[76]: cat_df['policy_cs1'].unique()
```

```
[76]: array(['250/500', '100/300', '500/1000'], dtype=object)
```

```
[77]: cat_df['insured_sex'].unique()
```

```
[77]: array(['MALE', 'FEMALE'], dtype=object)
```

```
[78]: cat_df['insured_education_level'].unique()
```

```
[78]: array(['MD', 'PhD', 'Associate', 'Masters', 'High School', 'College',
          'JD'], dtype=object)
```

```
[79]: cat_df['insured_relationship'].unique()
```

```
[79]: array(['husband', 'other-relative', 'own-child', 'unmarried', 'wife',
          'not-in-family'], dtype=object)
```

```
[80]: cat_df['incident_type'].unique()
```

```
[80]: array(['Single Vehicle Collision', 'Vehicle Theft',
          'Multi-vehicle Collision', 'Parked Car'], dtype=object)
```

```
[81]: cat_df['collision_type'].unique()
```

```
[81]: array(['Side Collision', 'Rear Collision', 'Front Collision'],
          dtype=object)
```

```
[82]: cat_df['incident_severity'].unique()
```

```
[82]: array(['Major Damage', 'Minor Damage', 'Total Loss', 'Trivial Damage'],
      dtype=object)
```

```
[83]: cat_df['authorities_contacted'].unique()
```

```
[83]: array(['Police', nan, 'Fire', 'Other', 'Ambulance'], dtype=object)
```

```
[84]: cat_df['property_damage'].unique()
```

```
[84]: array(['YES', 'NO'], dtype=object)
```

```
[85]: cat_df['police_report_available'].unique()
```

```
[85]: array(['YES', 'NO'], dtype=object)
```

```
[86]: cat_df['fraud_reported'].unique()
```

```
[86]: array(['Y', 'N'], dtype=object)
```

```
[87]: # custom mapping for encoding
cat_df['policy_cs1'] = cat_df['policy_cs1'].map({'100/300' : 1, '250/500' : 2.5,
↪ '500/1000' : 5})
cat_df['insured_education_level'] = cat_df['insured_education_level'].map({'JD' :
↪ 1, 'High School' : 2, 'College' : 3, 'Masters' : 4, 'Associate' : 5, 'MD' : 6, 'PhD' : 7})
cat_df['incident_severity'] = cat_df['incident_severity'].map({'Trivial Damage' :
↪ 1, 'Minor Damage' : 2, 'Major Damage' : 3, 'Total Loss' : 4})
cat_df['insured_sex'] = cat_df['insured_sex'].map({'FEMALE' : 0, 'MALE' : 1})
cat_df['property_damage'] = cat_df['property_damage'].map({'NO' : 0, 'YES' : 1})
cat_df['police_report_available'] = cat_df['police_report_available'].map({'NO' :
↪ 0, 'YES' : 1})
cat_df['fraud_reported'] = cat_df['fraud_reported'].map({'N' : 0, 'Y' : 1})
```

```
[88]: # auto encoding of categorical variables
for col in cat_df:
    ↪ drop(columns=['policy_cs1', 'insured_education_level', 'incident_severity', 'insured_sex', 'prop
    ↪ columns:
        cat_df= pd.get_dummies(cat_df, columns=[col], prefix = [col],
    ↪ drop_first=True)
```

```
[89]: # Converting any True/False values to 1/0 to make all data numeric
cat_df = cat_df.applymap(lambda x: int(x) if isinstance(x, bool) else x)
```

```
C:\Users\hamid\AppData\Local\Temp\ipykernel_29924\720693957.py:2: FutureWarning:
DataFrame.applymap has been deprecated. Use DataFrame.map instead.
    cat_df = cat_df.applymap(lambda x: int(x) if isinstance(x, bool) else x)
```

```
[90]: # data after encoding
cat_df.head()
```

```

[90]: policy_cs1  insured_sex  insured_education_level  incident_severity  \
0          2.5          1          6          3
1          2.5          1          6          2
2          1.0          0          7          2
3          2.5          0          7          3
4          5.0          1          5          2

    property_damage  police_report_available  fraud_reported  \
0                1                1                1
1                0                0                1
2                0                0                0
3                0                0                1
4                0                0                0

    insured_occupation_armed-forces  insured_occupation_craft-repair  \
0                                0                                1
1                                0                                0
2                                0                                0
3                                1                                0
4                                0                                0

    insured_occupation_exec-managerial  ...  insured_relationship_unmarried  \
0                                0  ...                                0
1                                0  ...                                0
2                                0  ...                                0
3                                0  ...                                1
4                                0  ...                                1

    insured_relationship_wife  incident_type_Parked Car  \
0                0                0
1                0                0
2                0                0
3                0                0
4                0                0

    incident_type_Single Vehicle Collision  incident_type_Vehicle Theft  \
0                1                0
1                0                1
2                0                0
3                1                0
4                0                1

    collision_type_Rear Collision  collision_type_Side Collision  \
0                0                1
1                1                0
2                1                0
3                0                0

```

```

4                                1                                0

    authorities_contacted_Fire  authorities_contacted_Other  \
0                                0                                0
1                                0                                0
2                                0                                0
3                                0                                0
4                                0                                0

    authorities_contacted_Police
0                                1
1                                1
2                                1
3                                1
4                                0

```

[5 rows x 33 columns]

```

[91]: # data after encoding
cat_df.head()

```

```

[91]:  policy_cs1  insured_sex  insured_education_level  incident_severity  \
0          2.5          1          6          3
1          2.5          1          6          2
2          1.0          0          7          2
3          2.5          0          7          3
4          5.0          1          5          2

    property_damage  police_report_available  fraud_reported  \
0                1                1                1
1                0                0                1
2                0                0                0
3                0                0                1
4                0                0                0

    insured_occupation_armed-forces  insured_occupation_craft-repair  \
0                                0                                1
1                                0                                0
2                                0                                0
3                                1                                0
4                                0                                0

    insured_occupation_exec-managerial  ...  insured_relationship_unmarried  \
0                                0  ...                                0
1                                0  ...                                0
2                                0  ...                                0
3                                0  ...                                1

```

```

4                                0 ...                                1

insured_relationship_wife    incident_type_Parked Car \
0                             0                             0
1                             0                             0
2                             0                             0
3                             0                             0
4                             0                             0

incident_type_Single Vehicle Collision    incident_type_Vehicle Theft \
0                                         1                             0
1                                         0                             1
2                                         0                             0
3                                         1                             0
4                                         0                             1

collision_type_Rear Collision    collision_type_Side Collision \
0                             0                             1
1                             1                             0
2                             1                             0
3                             0                             0
4                             1                             0

authorities_contacted_Fire    authorities_contacted_Other \
0                             0                             0
1                             0                             0
2                             0                             0
3                             0                             0
4                             0                             0

authorities_contacted_Police
0                             1
1                             1
2                             1
3                             1
4                             0

[5 rows x 33 columns]

```

```
[92]: # extracting the numerical columns
num_df = data.select_dtypes(include=['int64']).copy()
```

```
[93]: num_df.columns
```

```
[93]: Index(['months_as_customer', 'policy_deductable', 'umbrella_limit',
            'capital-gains', 'capital-loss', 'incident_hour_of_the_day',
            'number_of_vehicles_involved', 'bodily_injuries', 'witnesses',
```



```
    'injury_claim', 'property_claim', 'vehicle_claim'],
    dtype='object')
```

```
[94]: num_df.head()
```

```
[94]:   months_as_customer  policy_deductable  umbrella_limit  capital-gains  \
0                328                1000                0          53300
1                228                2000            5000000                0
2                134                2000            5000000          35100
3                256                2000            6000000          48900
4                228                1000            6000000          66000

   capital-loss  incident_hour_of_the_day  number_of_vehicles_involved  \
0              0                      5                      1
1              0                      8                      1
2              0                      7                      3
3          -62400                      5                      1
4          -46000                     20                      1

   bodily_injuries  witnesses  injury_claim  property_claim  vehicle_claim
0                1          2          6510          13020          52080
1                0          0           780           780           3510
2                2          3          7700          3850          23100
3                1          2          6340          6340          50720
4                0          1          1300           650           4550
```

```
[95]: # combining the Numerical and categorical dataframes to get the final dataset
final_df=pd.concat([num_df,cat_df], axis=1)
```

```
[96]: final_df.head()
```

```
[96]:   months_as_customer  policy_deductable  umbrella_limit  capital-gains  \
0                328                1000                0          53300
1                228                2000            5000000                0
2                134                2000            5000000          35100
3                256                2000            6000000          48900
4                228                1000            6000000          66000

   capital-loss  incident_hour_of_the_day  number_of_vehicles_involved  \
0              0                      5                      1
1              0                      8                      1
2              0                      7                      3
3          -62400                      5                      1
4          -46000                     20                      1

   bodily_injuries  witnesses  injury_claim  ...  \
0                1          2          6510  ...
```

1	0	0	780	...
2	2	3	7700	...
3	1	2	6340	...
4	0	1	1300	...

	insured_relationship_unmarried	insured_relationship_wife	\
0	0	0	
1	0	0	
2	0	0	
3	1	0	
4	1	0	

	incident_type_Parked Car	incident_type_Single Vehicle Collision	\
0	0	1	
1	0	0	
2	0	0	
3	0	1	
4	0	0	

	incident_type_Vehicle Theft	collision_type_Rear Collision	\
0	0	0	
1	1	1	
2	0	1	
3	0	0	
4	1	1	

	collision_type_Side Collision	authorities_contacted_Fire	\
0	1	0	
1	0	0	
2	0	0	
3	0	0	
4	0	0	

	authorities_contacted_Other	authorities_contacted_Police
0	0	1
1	0	1
2	0	1
3	0	1
4	0	0

[5 rows x 45 columns]

[]:

[97]: *# Checking for outliers in numerical columns*

```
import matplotlib.pyplot as plt
```

```

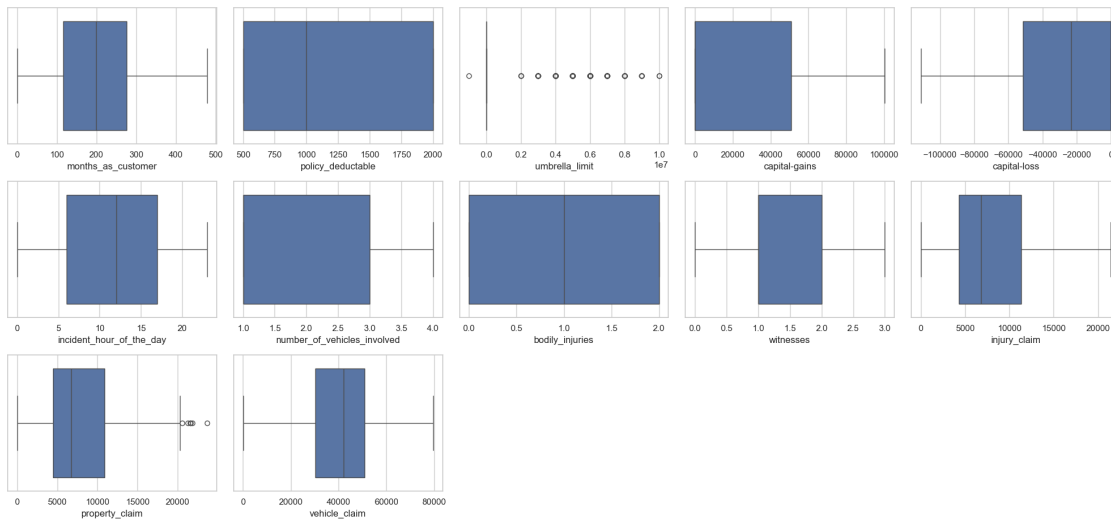
import seaborn as sns

plt.figure(figsize=(20, 15))
plot_number = 1

for col in num_df.columns:
    if plot_number <= 24: # Limiting to 24 plots for readability
        ax = plt.subplot(5, 5, plot_number) # Creating a 5x5 grid of boxplots
        sns.boxplot(x=num_df[col], ax=ax)
        plt.xlabel(col, fontsize=12)
        plot_number += 1

plt.tight_layout()
plt.show()

```



```

[98]: # Checking for potential outliers using IQR, without removing them yet
outliers_summary = {}

for col in num_df:
    Q1 = final_df[col].quantile(0.25)
    Q3 = final_df[col].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR
    outliers_count = final_df[(final_df[col] < lower_bound) | (final_df[col] >
    ↪upper_bound)].shape[0]

# Storing summary information
outliers_summary[col] = {

```

```

        "Lower Bound": lower_bound,
        "Upper Bound": upper_bound,
        "Outliers Count": outliers_count
    }

# Displaying the outliers summary
pd.DataFrame(outliers_summary).T

```

```

[98]:

```

	Lower Bound	Upper Bound	Outliers Count
months_as_customer	-125.0	517.0	0.0
policy_deductable	-1750.0	4250.0	0.0
umbrella_limit	0.0	0.0	202.0
capital-gains	-76537.5	127562.5	0.0
capital-loss	-128750.0	77250.0	0.0
incident_hour_of_the_day	-10.5	33.5	0.0
number_of_vehicles_involved	-2.0	6.0	0.0
bodily_injuries	-3.0	5.0	0.0
witnesses	-0.5	3.5	0.0
injury_claim	-6220.0	21820.0	0.0
property_claim	-5215.0	20545.0	6.0
vehicle_claim	-502.5	81617.5	0.0

```

[99]: # Cap outliers based on IQR bounds
for col in ['umbrella_limit', 'property_claim', 'vehicle_claim']:
    Q1 = final_df[col].quantile(0.25)
    Q3 = final_df[col].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR

    # Apply capping
    final_df[col] = np.where(final_df[col] < lower_bound, lower_bound,
↪final_df[col])
    final_df[col] = np.where(final_df[col] > upper_bound, upper_bound,
↪final_df[col])

print("Outliers have been capped to IQR bounds.")

```

Outliers have been capped to IQR bounds.

```

[100]: # Checking for outliers again after capping
outliers_summary = {}

for col in ['umbrella_limit', 'property_claim', 'vehicle_claim']:
    Q1 = final_df[col].quantile(0.25)
    Q3 = final_df[col].quantile(0.75)
    IQR = Q3 - Q1

```

```

lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR
outliers_count = final_df[(final_df[col] < lower_bound) | (final_df[col] >
↳upper_bound)].shape[0]

# Storing summary information
outliers_summary[col] = {
    "Lower Bound": lower_bound,
    "Upper Bound": upper_bound,
    "Outliers Count": outliers_count
}

# Displaying the outliers summary after capping
pd.DataFrame(outliers_summary).T

```

```

[100]:
           Lower Bound  Upper Bound  Outliers Count
umbrella_limit         0.0          0.0           0.0
property_claim       -5215.0       20545.0           0.0
vehicle_claim        -502.5       81617.5           0.0

```

```

[101]: # Define num_df to include only numerical columns
num_df = final_df.select_dtypes(include=[np.number]).columns

```

```

[102]: # Apply Min-Max Scaling
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
final_df[num_df] = scaler.fit_transform(final_df[num_df])

print("Feature scaling (Min-Max) applied to numerical columns.")

```

Feature scaling (Min-Max) applied to numerical columns.

```

[103]: # Check min and max after scaling
print(final_df[num_df].describe().loc[['min', 'max']])

```

```

      months_as_customer  policy_deductable  umbrella_limit  capital-gains  \
min                    0.0                0.0             0.0            0.0
max                    1.0                1.0             0.0            1.0

      capital-loss  incident_hour_of_the_day  number_of_vehicles_involved  \
min                0.0                    0.0                        0.0
max                1.0                    1.0                        1.0

      bodily_injuries  witnesses  injury_claim  ...  \
min                0.0          0.0          0.0  ...
max                1.0          1.0          1.0  ...

      insured_relationship_unmarried  insured_relationship_wife  \

```

```

min                0.0                0.0
max                1.0                1.0

incident_type_Parked Car  incident_type_Single Vehicle Collision \
min                0.0                0.0
max                1.0                1.0

incident_type_Vehicle Theft  collision_type_Rear Collision \
min                0.0                0.0
max                1.0                1.0

collision_type_Side Collision  authorities_contacted_Fire \
min                0.0                0.0
max                1.0                1.0

authorities_contacted_Other  authorities_contacted_Police
min                0.0                0.0
max                1.0                1.0

[2 rows x 45 columns]

```

Summary: Ready for Export and Next Steps

Alright, our data is now fully prepped! We've taken care of:

- **Handling missing values:** Replaced missing values to ensure consistency.
- **Encoding categorical variables:** Converted categorical data to numerical format so models can work with it.
- **Outlier treatment:** Capped extreme values to keep our data balanced.
- **Feature scaling:** Standardized the numerical columns to a [0, 1] range for smoother model performance.

Now we're ready to export this cleaned and processed DataFrame to a CSV file. This exported file will be our finalized dataset, and here's what we can do with it next:

1. **Training and Testing:** Use this file to split the data into training and test sets for building and evaluating our machine learning models.
2. **Deployment:** Since the data is clean, standardized, and model-ready, we can also use this same file in deployment for real-world predictions.

```
[104]: data.head()
```

```

[104]:  months_as_customer  policy_csl  policy_deductable  policy_annual_premium \
0                328      250/500                1000          1406.91
1                228      250/500                2000          1197.22
2                134      100/300                2000          1413.14
3                256      250/500                2000          1415.74
4                228      500/1000               1000          1583.91

umbrella_limit  insured_sex  insured_education_level  insured_occupation \

```

0	0	MALE		MD	craft-repair
1	5000000	MALE		MD	machine-op-inspct
2	5000000	FEMALE		PhD	sales
3	6000000	FEMALE		PhD	armed-forces
4	6000000	MALE		Associate	sales

	insured_relationship	capital-gains	...	incident_hour_of_the_day	\
0	husband	53300	...	5	
1	other-relative	0	...	8	
2	own-child	35100	...	7	
3	unmarried	48900	...	5	
4	unmarried	66000	...	20	

	number_of_vehicles_involved	property_damage	bodily_injuries	witnesses	\
0	1	YES	1	2	
1	1	NO	0	0	
2	3	NO	2	3	
3	1	NO	1	2	
4	1	NO	0	1	

	police_report_available	injury_claim	property_claim	vehicle_claim	\
0	YES	6510	13020	52080	
1	NO	780	780	3510	
2	NO	7700	3850	23100	
3	NO	6340	6340	50720	
4	NO	1300	650	4550	

	fraud_reported
0	Y
1	Y
2	N
3	Y
4	N

[5 rows x 25 columns]

```
[105]: # Exporting the final cleaned and preprocessed data to a CSV file
final_df.to_csv('insuranceFraud_final_processed_data.csv', index=False)
print("Data exported to 'insuranceFraud_final_processed_data.csv' successfully!")
↵")
```

Data exported to 'insuranceFraud_final_processed_data.csv' successfully!

```
[ ]:
```